

## Ex.No 1: Procedure to set up single node cluster using hadoop


### Step 1: Verify the Java installed

javac -version

### Step 2: Extract Hadoop at C:\Hadoop


### Step 3: Setting up the HADOOP\_HOME variable

Use windows environment variable setting for Hadoop Path setting.

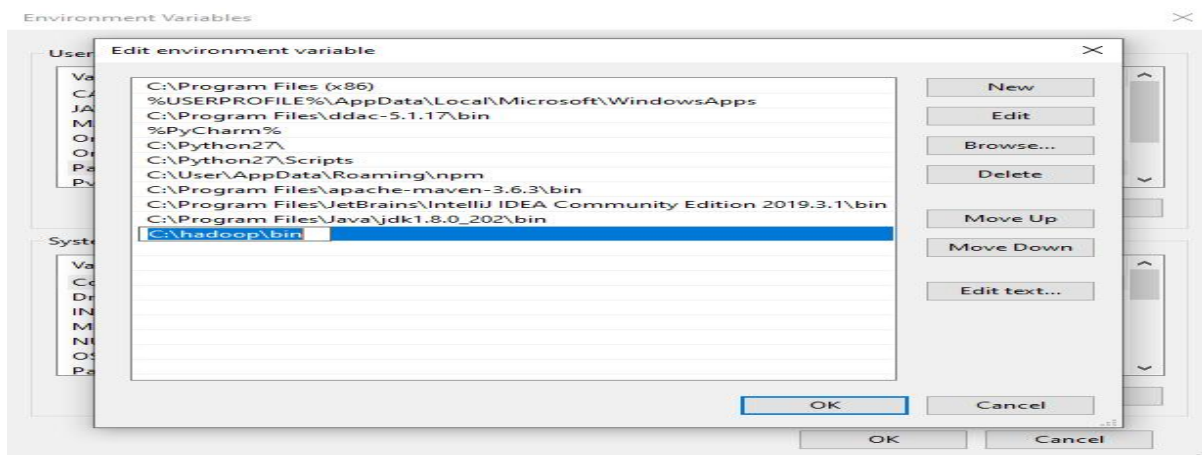


### Step 4: Set JAVA\_HOME variable

Use windows environment variable setting for Hadoop Path setting.



### Step 5: Set Hadoop and Java bin directory path



### Step 6: Hadoop Configuration :

For Hadoop Configuration we need to modify Six files that are listed below-

1. Core-site.xml

2. Mapred-site.xml
3. Hdfs-site.xml
4. Yarn-site.xml
5. Hadoop-env.cmd
6. Create two folders datanode and namenode

### **Step 6.1: Core-site.xml configuration**

```
<configuration>
  <property>
    <name>fs.defaultFS</name>
    <value>hdfs://localhost:9000</value>
  </property>
</configuration>
```

### **Step 6.2: Mapred-site.xml configuration**

```
<configuration>
  <property>
    <name>mapreduce.framework.name</name>
    <value>yarn</value>
  </property>
</configuration>
```

### **Step 6.3: Hdfs-site.xml configuration**

```
<configuration>
  <property>
    <name>dfs.replication</name>
    <value>1</value>
  </property>
  <property>
    <name>dfs.namenode.name.dir</name>
    <value>C:\hadoop-2.8.0\data\namenode</value>
  </property>
  <property>
```

```

    <name>dfs.datanode.data.dir</name>
    <value>C:\hadoop-2.8.0\data\datanode</value>
  </property>
</configuration>

```

#### Step 6.4: Yarn-site.xml configuration

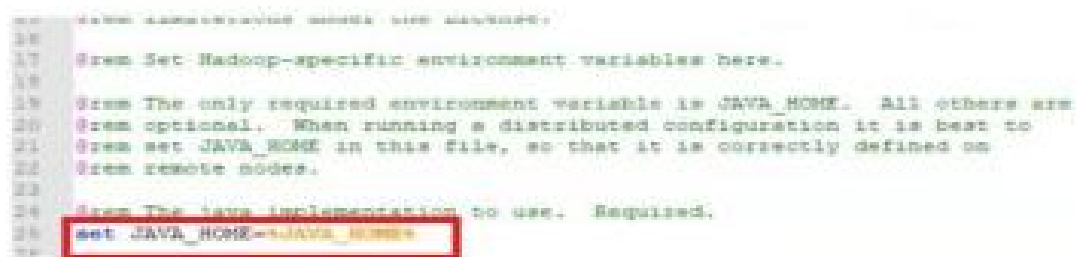
```

<configuration>
  <property>
    <name>yarn.nodemanager.aux-services</name>
    <value>mapreduce_shuffle</value>
  </property>
  <property>
    <name>yarn.nodemanager.auxservices.mapreduce.shuffle.class</name>
    <value>org.apache.hadoop.mapred.ShuffleHandler</value>
  </property>
</configuration>

```

#### Step 6.5: Hadoop-env.cmd configuration

Set "JAVA\_HOME=C:\Java" (On C:\java this is path to file jdk.18.0)



```

16  %rem Set Hadoop-specific environment variables here.
17
18  %rem The only required environment variable is JAVA_HOME. All others are
19  %rem optional. When running a distributed configuration it is best to
20  %rem set JAVA_HOME in this file, so that it is correctly defined on
21  %rem remote nodes.
22
23  %rem The java implementation to use. Required.
24  set JAVA_HOME=%JAVA_HOME%
25
26

```


#### Step 6.6: Create datanode and namenode folders

1. Create folder "data" under "C:\Hadoop-2.8.0"
2. Create folder "datanode" under "C:\Hadoop-2.8.0\data"
3. Create folder "namenode" under "C:\Hadoop-2.8.0\data"

#### Step 7: Format the namenode folder

Open command window (cmd) and typing command "hdfs namenode -format"

```
C:\hadoop\hadoop-2.7.1\sbin>start-all.cmd
```

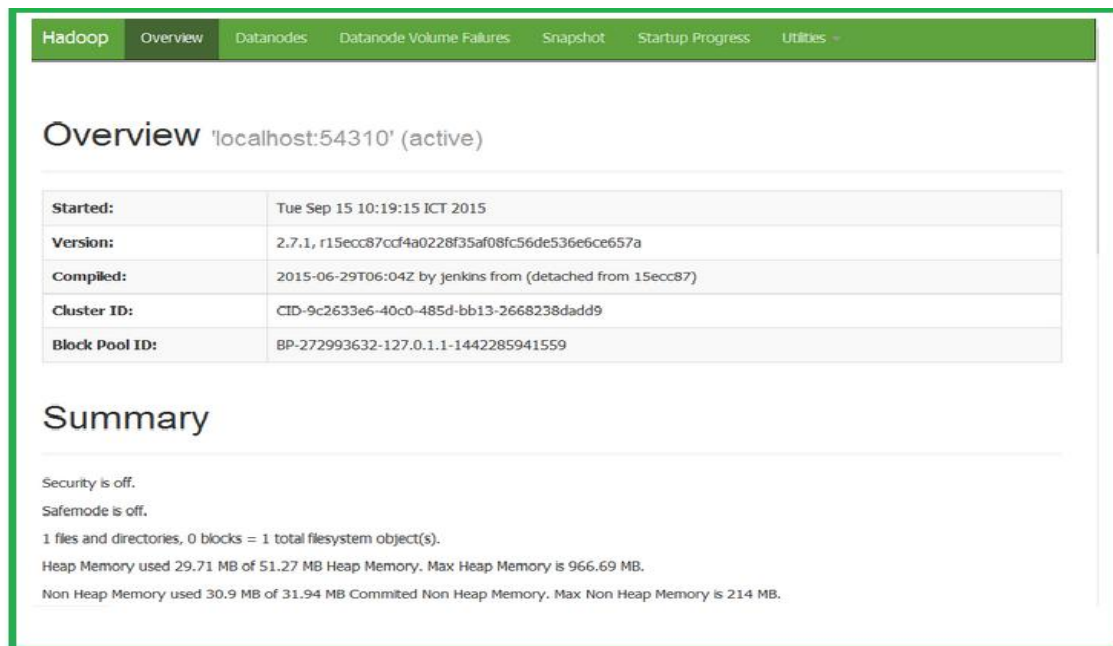


**Step 9:** Open: <http://localhost:8088>



The image shows the Hadoop All Applications page. On the left is a sidebar with a 'Cluster' section containing links for About, Nodes, Node Labels, Applications, and a 'Tools' section. The main content area is titled 'All Applications' and displays various metrics. At the top, 'Cluster Metrics' shows counts for Apps Submitted (1), Apps Pending (0), Apps Running (0), Apps Completed (0), Containers Running (0), Memory Used (0.0), and Memory Free (0.0). Below this, 'Cluster Nodes Metrics' shows Active Nodes (1), Decommissioning Nodes (0), Decommissioned Nodes (0), and Lost Nodes (1). The 'Scheduler Metrics' section shows Capacity Scheduler as the Scheduler Type and Capacity as the Scheduling Resource Type. A table at the bottom lists application details with columns: ID, User, Name, Application Type, Queue, Application Priority, StartTime, FinishTime, State, FinalState, Running Containers, Allocated CPU, and Allocated Memory. The first row shows an application with ID 1, User root, Name test, Application Type test, Queue test, Application Priority 1, StartTime 2015-09-15 10:19:15, FinishTime 2015-09-15 10:19:15, State Running, FinalState Running, 1 Running Containers, 100% Allocated CPU, and 100 MB Allocated Memory.

**Step 10:** Open: <http://localhost:50070>



The image shows the Hadoop Overview page. The top navigation bar includes links for Hadoop, Overview, Datanodes, Datanode Volume Failures, Snapshot, Startup Progress, and Utilities. The main content area is titled 'Overview' and shows the cluster is 'localhost:54310' (active). Below this is a table with the following information:

Started:	Tue Sep 15 10:19:15 ICT 2015
Version:	2.7.1, r15ecc87cd4a0228f35af08fc56de536e6ce657a
Compiled:	2015-06-29T06:04Z by jenkins from (detached from 15ecc87)
Cluster ID:	CID-9c2633e6-40c0-485d-bb13-2668238dadd9
Block Pool ID:	BP-272993632-127.0.1.1-1442285941559

Below the table is a 'Summary' section with the following text:

Security is off.  
Safemode is off.  
1 files and directories, 0 blocks = 1 total filesystem object(s).  
Heap Memory used 29.71 MB of 51.27 MB Heap Memory. Max Heap Memory is 966.69 MB.  
Non Heap Memory used 30.9 MB of 31.94 MB Committed Non Heap Memory. Max Non Heap Memory is 214 MB.

## Ex.No 2: Word Count Program to use of Map and Reduce Tasks

```
import java.io.IOException;
import java.util.StringTokenizer;
import org.apache.hadoop.conf.Configuration;
import org.apache.hadoop.fs.Path;
import org.apache.hadoop.io.IntWritable;
import org.apache.hadoop.io.Text;
import org.apache.hadoop.mapreduce.Job;
import org.apache.hadoop.mapreduce.Mapper;
import org.apache.hadoop.mapreduce.Reducer;
import org.apache.hadoop.mapreduce.lib.input.FileInputFormat;
import org.apache.hadoop.mapreduce.lib.output.FileOutputFormat;
public class WordCount {
    public static class TokenizerMapper
        extends Mapper<Object, Text, Text, IntWritable>{
        private final static IntWritable one = new IntWritable(1);
        private Text word = new Text();
        public void map(Object key, Text value, Context context
            ) throws IOException, InterruptedException {
            StringTokenizer itr = new StringTokenizer(value.toString());
            while (itr.hasMoreTokens()) {
                word.set(itr.nextToken());
                context.write(word, one);
            }
        }
    }

    public static class IntSumReducer
        extends Reducer<Text,IntWritable,Text,IntWritable> {
        private IntWritable result = new IntWritable();
        public void reduce(Text key, Iterable<IntWritable> values,
            Context context
            ) throws IOException, InterruptedException {
            int sum = 0;
            for (IntWritable val : values) {
                sum += val.get();
            }
            result.set(sum);
            context.write(key, result);
        }
    }

    public static void main(String[] args) throws Exception {
```

```
Configuration conf = new Configuration();
Job job = Job.getInstance(conf, "word count");
job.setJarByClass(WordCount.class);
job.setMapperClass(TokenizerMapper.class);
job.setCombinerClass(IntSumReducer.class);
job.setReducerClass(IntSumReducer.class);
job.setOutputKeyClass(Text.class);
job.setOutputValueClass(IntWritable.class);
FileInputFormat.addInputPath(job, new Path(args[0]));
FileOutputFormat.setOutputPath(job, new Path(args[1]));
System.exit(job.waitForCompletion(true) ? 0 : 1);
}
}
```

**Output:**

```
C:\hadoop\hadoop-2.7.1\bin>set classpath=c:\hadoop\hadoop-core-1.2.1.jar
C:\hadoop\hadoop-2.7.1\bin>set classpath=c:\program files\java\jdk1.7.0\bin
C:\hadoop\hadoop-2.7.1\bin>javac WordCount.java
C:\hadoop\hadoop-2.7.1\bin>jar -cvf wc.jar WordCount*.class
C:\hadoop\hadoop-2.7.1\bin>hadoop fs -mkdir /f1
C:\hadoop\hadoop-2.7.1\bin>hadoop fs -cat file1.txt
Hello world Bye world
C:\hadoop\hadoop-2.7.1\bin>hadoop fs -put file1.txt /f1
C:\hadoop\hadoop-2.7.1\bin>hadoop jar wc.jar WordCount /f1 /f2
```

**File System Counters**

```
FILE: Number of bytes read=6604
FILE: Number of bytes written=566732
FILE: Number of read operations=0
FILE: Number of large read operations=0
FILE: Number of write operations=0
HDFS: Number of bytes read=42
HDFS: Number of bytes written=22
HDFS: Number of read operations=13
HDFS: Number of large read operations=0
HDFS: Number of write operations=4
```

**Map-Reduce Framework**

```
Map input records=1
Map output records=4
Map output bytes=38
Map output materialized bytes=40
Input split bytes=99
Combine input records=4
```



Combine output records=3

Reduce input groups=3

Reduce shuffle bytes=40

Reduce input records=3

Reduce output records=3

Spilled Records=6

Shuffled Maps =1

Failed Shuffles=0

Merged Map outputs=1

GC time elapsed (ms)=0

Total committed heap usage (bytes)=641728512

#### Shuffle Errors

BAD\_ID=0

CONNECTION=0

IO\_ERROR=0

WRONG\_LENGTH=0

WRONG\_MAP=0

WRONG\_REDUCE=0

#### File Input Format Counters

Bytes Read=21

#### File Output Format Counters

Bytes Written=22

C:\hadoop\hadoop-2.7.1\sbin>hadoop fs -cat /f3/part-r-00000

Bye 1

Hello 1

world 2

### Ex.No 3: Python basic for pandas

1. import pandas as pd
2. # Reading data from google drive in to a data frame  
data=pd.read\_csv("/content/drive/My Drive/ Code/data/cafe.csv")
3. # Viewing data head (default: first 5 rows)  
data.head()
4. # Sum of a column  
data["AFF-Mon"].sum()
5. # Minimum value in a column  
data["AFF-Mon"].min()
6. # Numbr of values in a column  
data["AFF-Mon"].count()
7. # Mean value of a column  
data["AFF-Mon"].mean()
8. # Basic statistical values of a column  
data["AFF-Mon"].describe()
9. # STatistics of the whole data frame  
data.describe()
10. # STatistics of the whole data frame including all columns (numeric and non-numeric)  
data.describe(include="all")
11. # View last 5 rows of a data frame  
data.tail()
12. # View column names  
data.columns
13. # Unique values in a column  
data["State"].unique()
14. # Number of unique values in a column  
data["State"].nunique()
15. # Number of rows and columns in a data frame  
data.shape

```

16. # Add column to a data frame
    # First create a list of numbers for 85 rows
    x=[i for i in range(85)]
    # Assign the list to new column
    data["New Index"]=x

17. # View column names
    data.columns

18. data.shape

19. data.head()

20. # Dropping (removing) a column
    data.drop("New Index",axis=1,inplace=True) # axis=1 for columns and axis=0 for rows.

21. # Conditional statements (or filtering data using conditions)
22. # Example, all store data from the state Tamil Nadu (TN)
    data[data["State"]=="TN"]

23. # Detecting missing values
    data.isna()

24. # Detecting missing values in a single column
    data['Store ID'].isna()

25. # Number of missing values per column
    data.isna().sum()

26. # Number of missing values in the data frame
    data.isna().sum().sum()

27. # Reading "train.csv" data from google drive in to a data frame
    df=pd.read_csv("/content/drive/My Drive/Colab Notebooks/ATAL FDP GGV Python
    Code/data/train.csv")

28. # View first 5 rows
    df.head()

29. # Detecting missing values
    df.isna().sum()

30. # Filling missing values with mean data
31. df["Age"]=df["Age"].fillna(df["Age"].mean())

32. df.isna().sum()

```

## Ex.No 4: Python Code for Matplotlib

1. `import matplotlib.pyplot as plt`
2. `plt.plot([1,2,3,4,5],[1,4,6,8,9],"ro")`  
`plt.show()`
3. `plt.plot([1,2,3,4,5],[1,4,6,8,9])`  
`plt.title("Graph")`  
`plt.xlabel("X-axis")`  
`plt.ylabel("Y-axis")`  
`plt.show()`
4. `import numpy as np`
5. `t=np.arange(0.,5.,0.2)`  
`t`
6. `array([0. , 0.2, 0.4, 0.6, 0.8, 1. , 1.2, 1.4, 1.6, 1.8, 2. , 2.2, 2.4,`  
`2.6, 2.8, 3. , 3.2, 3.4, 3.6, 3.8, 4. , 4.2, 4.4, 4.6, 4.8])`  
`plt.plot(t,t,"r--",t,t**2,'bs',t,t**3,'g^')`  
`plt.title("Linear, Quadratic and Cubic graphs")`  
`plt.xlabel("X-axis")`  
`plt.ylabel("Y-axis")`  
`plt.show()`
7. `x1=[5,8,10]`  
`y1=[12,6,16]`  
`x2=[6,9,11]`  
`y2=[6,15,17]`  
`plt.plot(x1,y1,"green",label="Age",linewidth=5)`  
`plt.plot(x2,y2,"blue",label="Income",linewidth=5)`  
`plt.legend()`  
`plt.grid(True,color='lightgrey')`  
`plt.show()`
8. `# Bar graph`  
`x1=[5,8,10]`  
`y1=[12,6,16]`  
`plt.bar(x1,y1,color="green")`  
`plt.show()`
9. `# Scatter plot`  
`x=[6,9,11,14,20,24,30]`  
`y=[6,15,17,22,30,37,44]`  
`plt.scatter(x,y)`

```
plt.show()
```

```
10. # Plotting categorial graph
```

```
x=[100,200,300,400]
```

```
y=[12,20,30,40]
```

```
plt.figure(figsize=(9,3))
```

```
plt.subplot(131)
```

```
plt.bar(x,y)
```

```
plt.subplot(132)
```

```
plt.scatter(x,y)
```

```
plt.subplot(133)
```

```
plt.plot(x,y)
```

```
11. import pandas as pd
```

```
data=pd.read_csv("/content/drive/My Drive/cafe.csv")
```

```
data.head()
```

```
12. addr=data["Address"].head()
```

```
addr
```

```
13. mon=data["AFF-Mon"].head()
```

```
plt.bar(addr,mon)
```

```
plt.show()
```

```
14. plt.bar(addr,mon)
```

```
plt.xticks(addr,('a','b','c','d','e'))
```

```
plt.shw()
```

```
15. tue=data["AFF-Tue"].head()
```

```
wed=data["AFF-Wed"].head()
```

```
# Saving graph to a file
```

```
plt.plot(addr,mon,'r--',addr,tue,'bs',addr,wed,'g^')
```

```
plt.xticks(rotation=90)
```

```
plt.savefig("/content/drive/My Drive/Colab Notebooks/MyGraph.png",format='png')
```

```
plt.show()
```

```
16. plt.scatter(addr,mon)
```

```
17. plt.xticks(rotation=90)
```

```
18. plt.show()
```

```
19. import folium as f
```

```
f.Map(location=[18.932308,72.834091])
```

## Ex.No 5: Python Code for Seaborn

1. 

```
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
data=pd.read_csv("/content/drive/My Drive/train.csv")
data.head()
```
2. 

```
# Default Seaborn datasets
sns.get_dataset_names()
```
3. 

```
# Loading Seaborn predefined dataset
df=sns.load_dataset('tips')
df.head()
```
4. 

```
# Range and outliers
plt.scatter(data["Sex"],data["Age"])
plt.show()
```
5. 

```
plt.style.use("seaborn")
fig,ax=plt.subplots(figsize=(15,9))
ax.boxplot([data["Fare"]])
plt.show()
```
6. 

```
data["Fare"].max()
512.3292
```
7. 

```
# Identifying outliers using z-score
z_score=(data["Age"]-data["Age"].mean())/data["Age"].std()
print(z_score)
```
8. 

```
# Example: use only data where z-score is less than 3
z_score<3
```
9. 

```
# Adding z_score column to data
data['z_score']=z_score
data.head()
```
10. 

```
# Filtering data where z_score < 3
data[data['z_score']<3]
```
11. 

```
data.skew()
```

```

12. # Histogram (Frequency distribution)
    sns.displot(data["Age"],kde=True)
    plt.show()

13. data.groupby(["Sex","Ticket"]).describe()

14. data.groupby(["Embarked","Sex","Pclass"]).size()
    data.groupby(["Embarked","Sex","Pclass"]).size().unstack()

15. # Arranging data into bins
    group_data=pd.DataFrame(pd.cut(data["Age"],bins=[0,15,30,45,60,75,90]))
    group_data.head()

16. group_data.groupby(["Age"]).size()

17. # Bar graph
    group_data.groupby(["Age"]).size().plot(kind="bar")
    plt.show()

18. # Pie plot
    group_data.groupby(["Age"]).size().plot(kind="pie")
    plt.show()

    # 3D plots
19. from mpl_toolkits import mplot3d
    fig=plt.figure()
    ax=plt.axes(projection="3d")
    plt.show()

20. ax=plt.axes(projection="3d")
    z=np.linspace(0,15,1000)
    x=np.sin(z)
    y=np.cos(z)
    ax.plot3D(x,y,z,"grey")

    zdata=15*np.random.random(100)
    xdata=np.sin(zdata)+0.1*np.random.random(100)
    ydata=np.cos(zdata)+0.1*np.random.random(100)

    plt.figure(figsize=(10,7))
    sns.pairplot(data)
    plt.show()

```

## Ex.No 6: Regression Algorithm

1. `import pandas as pd`
2. `Stock_Market = {'Year':  
[2017,2017,2017,2017,2017,2017,2017,2017,2017,2017,2017,2016,2016,2016,2016,  
2016,2016,2016,2016,2016,2016,2016,2016],  
'Month': [12, 11,10,9,8,7,6,5,4,3,2,1,12,11,10,9,8,7,6,5,4,3,2,1],  
'Interest_Rate':  
[2.75,2.5,2.5,2.5,2.5,2.5,2.5,2.25,2.25,2.25,2.2,2,1.75,1.75,1.75,1.75,1.75,1.75,1.75,1.  
75,1.75,1.75],  
'Unemployment_Rate':  
[5.3,5.3,5.3,5.3,5.4,5.6,5.5,5.5,5.5,5.5,5.6,5.7,5.9,6,5.9,5.8,6.1,6.2,6.1,6.1,6.1,5.9,6.2,6.2,6.1],  
'Stock_Index_Price':  
[1464,1394,1357,1293,1256,1254,1234,1195,1159,1167,1130,1075,1047,965,943,958,971,  
949,884,866,876,822,704,719]  
}`
3. `df = pd.DataFrame Stock_Market,columns=  
['Year','Month','Interest_Rate','Unemployment_Rate', 'Stock_Index_Price'])  
print (df)`
4. `import os  
os.chdir('/content/drive/My Drive/')  
os.getcwd()`
5. `df.to_csv('data/stock.csv', index=False)  
df=pd.read_csv('data/stock.csv')  
df.head()`
6. `# Check for Linearity  
import matplotlib.pyplot as plt  
plt.scatter(df['Interest_Rate'], df['Stock_Index_Price'], color='red')  
plt.title('Stock Index Price Vs Interest Rate', fontsize=14)  
plt.xlabel('Interest Rate', fontsize=14)  
plt.ylabel('Stock Index Price', fontsize=14)  
plt.grid(True)  
plt.show()`
7. `plt.scatter(df['Unemployment_Rate'], df['Stock_Index_Price'], color='blue')  
plt.title('Stock Index Price Vs Unemployment Rate', fontsize=14)  
plt.xlabel('Unemployment Rate', fontsize=14)  
plt.ylabel('Stock Index Price', fontsize=14)  
plt.grid(True)  
plt.show()`



```

8. from sklearn import linear_model
import statsmodels.api as sm

9. # Perform Multi-Linear Regression
# here we have 2 variables for multiple regression. If you just want to use one
# variable for simple linear regression, then use X = df['Interest_Rate'] for
# example. Alternatively, you may add additional variables within the brackets
X = df[['Interest_Rate','Unemployment_Rate']]
Y = df['Stock_Index_Price']
# with sklearn
regr = linear_model.LinearRegression()
regr.fit(X, Y)
print('Intercept: \n', regr.intercept_)
print('Coefficients: \n', regr.coef_)

# prediction with sklearn
New_Interest_Rate = 2.75
New_Unemployment_Rate = 5.3
print ('Predicted Stock Index Price: \n',
regr.predict([[New_Interest_Rate ,New_Unemployment_Rate]]))

# With StatModels
X = sm.add_constant(X) # adding a constant
model = sm.OLS(Y, X).fit() # Ordinary Least Squares
predictions = model.predict(X)
print_model = model.summary()
print(print_model)

```

## Ex.No 7: Model Training

1. `from google.colab import drive`  
`drive.mount('/content/drive')`
2. `import os`  
`os.chdir('/content/drive/My Drive/')`  
`os.getcwd()`
3. `from keras import layers`  
`from keras import models`
4. `model = models.Sequential()`
5. `model.add(layers.Conv2D(32, (3, 3), activation='relu',`  
`input_shape=(150, 150, 3)))`
6. `model.add(layers.MaxPooling2D((2, 2)))`  
`model.add(layers.Conv2D(64, (3, 3), activation='relu'))`  
`model.add(layers.MaxPooling2D((2, 2)))`  
`model.add(layers.Conv2D(128, (3, 3), activation='relu'))`  
`model.add(layers.MaxPooling2D((2, 2)))`  
`model.add(layers.Conv2D(128, (3, 3), activation='relu'))`  
`model.add(layers.MaxPooling2D((2, 2)))`  
`model.add(layers.Flatten())`  
`model.add(layers.Dropout(0.5))`  
`model.add(layers.Dense(512, activation='relu'))`  
`model.add(layers.Dense(1, activation='sigmoid'))`
7. `***Let's look at how the dimensions of the feature maps change with every successive`  
`layer:***`  
`model.summary()`
8. `# **Configuring the model for training**`  
`**For the compilation step, you'll go with the RMSprop optimizer. Because you ended the`  
`network with a single sigmoid unit, you'll use binary_crossentropy as the loss**`  
`from keras import optimizers`  
`model.compile(loss='binary_crossentropy',`  
`optimizer=optimizers.RMSprop(lr=1e-4),`  
`metrics=['accuracy'])`

## 9. # \*\*Data preprocessing\*\*

Currently, the data sits on a drive as JPEG files, so the steps for getting it into the network are roughly as follows:

1. Read the picture files.
2. Decode the JPEG content to RGB grids of pixels.
3. Convert these into floating-point tensors.
4. Rescale the pixel values (between 0 and 255) to the  $[0, 1]$  interval (as you know, neural networks prefer to deal with small input values).

**\*\*Using ImageDataGenerator to read images from directories\*\***

```
from keras.preprocessing.image import ImageDataGenerator
train_datagen = ImageDataGenerator(rescale=1./255)
#test_datagen = ImageDataGenerator(rescale=1./255)
train_generator = train_datagen.flow_from_directory('data/images',target_size=(150, 150),
batch_size=20, class_mode='binary')
```

10. The output of one of these generators: it yields batches of  $150 \times 150$  RGB images (shape (20, 150, 150, 3)) and binary labels (shape (20,)). There are 20 samples in each batch (the batch size).

```
for data_batch, labels_batch in train_generator:
    print('data batch shape:', data_batch.shape)
    print('labels batch shape:', labels_batch.shape)
    break
```

## 11. # \*\*Fitting the model using a batch generator\*\*

```
history = model.fit_generator( train_generator,steps_per_epoch=10,epochs=5)
```

## 12. # \*\*Saving the model\*\*

```
model.save('output/mymodel.h5')
```

## 13. ## \*\*Displaying curves of loss and accuracy during training\*\*

```
import matplotlib.pyplot as plt
acc = history.history['accuracy']
loss = history.history['loss']
epochs = range(1, len(acc) + 1)
plt.plot(epochs, acc, 'b', label='Training acc')
plt.title('Training accuracy')
plt.legend()
plt.figure()
plt.plot(epochs, loss, 'b', label='Training loss')
plt.title('Training loss')
plt.legend()
plt.show()
```

### **Ex.No 8: Neural\_Networks**

1. `import keras`  
`from keras.layers import Dense`  
`from keras.models import Sequential`
2. `import os`  
`os.chdir('/content/drive/My Drive/')`  
`os.getcwd()`
3. `import pandas as pd`  
`import numpy as np`
4. `ds=pd.read_csv("Iris.csv")`  
`ds.head()`
5. `arr=ds.values`  
`x=arr[:,0:4]`  
`y=arr[:,4]`
6. `model=Sequential()`  
`model.add(Dense(units=64,input_dim=4,activation="relu"))`  
`model.add(Dense(units=32,activation="relu"))`  
`model.add(Dense(units=3,activation="softmax"))`
7. `model.compile(loss='categorical_crossentropy', optimizer='adam',`  
`metrics=['accuracy'])`
8. `from sklearn.preprocessing import LabelEncoder`  
`from keras.utils import np_utils`  
`encoder = LabelEncoder()`  
`encoder.fit(y)`  
`encoded_Y = encoder.transform(y)`  
`dummy_y = np_utils.to_categorical(encoded_Y)`

9. `h=model.fit(x.astype('float'),dummy_y,epochs=10)`

Epoch 1/10

5/5 [=====] - 1s 5ms/step - loss: 3.8956 - accuracy:  
0.5867

Epoch 2/10

5/5 [=====] - 0s 3ms/step - loss: 1.5611 - accuracy:  
0.4000

Epoch 3/10

5/5 [=====] - 0s 4ms/step - loss: 0.9363 - accuracy:  
0.4800

Epoch 4/10

5/5 [=====] - 0s 3ms/step - loss: 0.9004 - accuracy:  
0.5667

Epoch 5/10

5/5 [=====] - 0s 4ms/step - loss: 0.7987 - accuracy:  
0.6133

Epoch 6/10

5/5 [=====] - 0s 3ms/step - loss: 0.6825 - accuracy:  
0.6600

Epoch 7/10

5/5 [=====] - 0s 3ms/step - loss: 0.6173 - accuracy:  
0.7933

Epoch 8/10

5/5 [=====] - 0s 5ms/step - loss: 0.6038 - accuracy:  
0.7667

Epoch 9/10

5/5 [=====] - 0s 3ms/step - loss: 0.5508 - accuracy:  
0.7733

Epoch 10/10

5/5 [=====] - 0s 3ms/step - loss: 0.5352 - accuracy:  
0.7867

```
10. from matplotlib import pyplot as plt
plt.plot(h.history['accuracy'])
plt.plot(h.history['loss'])
plt.ylabel('Accuracy/Loss')
plt.xlabel('Epoch')
plt.legend(['Accuracy', 'Loss'], loc='upper left')
plt.show()
```

